

Name: Zirui Zheng (Michael)
Course: Machine learning with R

Step 1 - Explore and Prepare the data

Explore the data by checking the structure and summary.

Structure:

```
'data.frame': 4521 obs. of 16 variables:
 $ age      : int  61 38 49 45 56 31 51 36 42 39 ...
 $ job      : Factor w/ 12 levels "admin.,"blue-collar",...: 11 3 2 5 10 10 8 7 5 10 ...
 $ marital  : Factor w/ 3 levels "divorced","married",...: 2 2 1 1 2 3 2 3 2 3 ...
 $ education: Factor w/ 4 levels "primary","secondary",...: 1 3 1 3 2 2 2 3 3 2 ...
 $ default  : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
 $ balance  : int  1414 4600 1584 1786 1176 2296 1840 1882 1958 1650 ...
 $ housing  : Factor w/ 2 levels "no","yes": 1 1 2 2 1 2 1 1 2 2 ...
 $ loan     : Factor w/ 2 levels "no","yes": 1 1 1 1 2 1 1 1 1 1 ...
 $ contact  : Factor w/ 3 levels "cellular","telephone",...: 1 3 3 3 1 3 2 1 1 1 ...
 $ day      : int  30 16 3 8 10 6 9 21 20 20 ...
 $ month    : Factor w/ 12 levels "apr","aug","dec",...: 5 7 7 9 6 7 4 2 10 2 ...
 $ campaign : int  2 4 2 2 1 2 4 3 1 2 ...
 $ pdays    : int  -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 ...
 $ previous : int  0 0 0 0 0 0 0 0 0 0 ...
 $ poutcome : Factor w/ 4 levels "failure","other",...: 4 4 4 4 4 4 4 4 4 4 ...
 $ y        : Factor w/ 2 levels "no","yes": 1 2 2 2 2 1 2 2 1 1 ...
```

Summary:

```
   age      job      marital  education  default  balance  housing  loan
Min. :21.00  management:969  divorced:528  primary :678  no :4445  Min. : -3499
1st Qu.:35.00 blue-collar:946  married :2797  secondary:2306  yes: 76  1st Qu.: 1529
Median:41.00 technician:768  single :1196  tertiary:1350      Median: 1900
Mean :43.17  admin. :478      unknown :187      Mean : 2869
3rd Qu.:51.00  services :417      3rd Qu.: 2925
Max. :89.00  retired :230      Max. :72751
(Other) :713
   contact  day      month  campaign  pdays  previous
poutcome
cellular:2896  Min. : 1.00  may :1398  Min. : 1.000  Min. : -1.00  Min. : 0.0000
failure: 490
telephone:301  1st Qu.: 9.00  jul :706  1st Qu.: 1.000  1st Qu.: -1.00  1st Qu.: 0.0000
other :197
unknown :1324  Median:16.00  aug :633  Median: 2.000  Median: -1.00  Median:
0.0000  success:129
      Mean :15.92  jun :531  Mean : 2.794  Mean :39.77  Mean : 0.5426
unknown:3705
      3rd Qu.:21.00  nov :389  3rd Qu.: 3.000  3rd Qu.: -1.00  3rd Qu.: 0.0000
      Max. :31.00  apr :293  Max. :50.000  Max. :871.00  Max. :25.0000
      (Other): 571
```

```
y
no :4000
yes: 521
```

Use any() function to see that there's no missing value:

```
[1] FALSE
```

After the exploration, I found that the data set is similar to the “credit.csv” file we did in class. There not not just numeric data in the file, and our predictor “y” is categorical by “yes” and “no”. Thus, I think it better to use a decision tree model to start: to build decision trees to decide if one subscribed a term deposit. As I am not going to use numerical methods like regression, thus, I am not going to use logistic algorithm, dummy variables, or to make it nominal distributed.

Step 2 - Train models

Sample the train data with 80% rate. After splitting the train and test data, check the proportion of y in both sets:

train:

```
      no      yes
0.8849558 0.1150442
```

test:

```
      no      yes
0.8839779 0.1160221
```

Use C5.0 function to build a decision tree model.

```
Call:
C5.0.default(x = f_train_dt[-16], y = f_train_dt$y)
```

```
Classification Tree
Number of samples: 3616
Number of predictors: 15
```

```
Tree size: 9
```

```
Non-standard options: attempt to group attributes
```

```
      no yes
867  38
```

Step 3 - Evaluate performance

The CrossTable() shows that kappa is 0.873.

Cell Contents

N
N / Row Total
N / Col Total
N / Table Total

Total Observations in Table: 905

	predicted y		
actual y	no	yes	Row Total
no	781	19	800
	0.976	0.024	0.884
	0.901	0.500	
	0.863	0.021	
yes	86	19	105
	0.819	0.181	0.116
	0.099	0.500	
	0.095	0.021	
Column Total	867	38	905
	0.958	0.042	

Error rate by checking the difference between prediction and test data:

[1] 0.1160221

Step 4 - Improve performance

**Use adabag library to do an adaboost on the model. However, the result doesn't look good in Kappa compared to decision tree above. Plus, I used detach() function to detach the adabag library, since I found it conflicting with ipred when using bagging method.

```

      Observed Class
Predicted Class no yes
no 3968 326
yes 32 195
Estimate Std.Err 2.5% 97.5% P-value
kappa 0.4854 0.0227 0.4409 0.5299 1.961e-101

```

**Use C5.0 to boost the decision tree. I ran a for loop so that I can check different numbers of trials (10, 20, 30). Also, in the for loop, I generated crosstable and mean error accordingly, followed by the increasing order of trials:

Cell Contents

--

	N
	N / Row Total
	N / Col Total
	N / Table Total

Total Observations in Table: 905

	predicted y		
actual y	no	yes	Row Total

no	785	15	800
	0.981	0.019	0.884
	0.895	0.536	
	0.867	0.017	

yes	92	13	105
	0.876	0.124	0.116
	0.105	0.464	
	0.102	0.014	

Column Total	877	28	905
	0.969	0.031	

Cell Contents

	N
	N / Row Total
	N / Col Total
	N / Table Total

Total Observations in Table: 905

	predicted y		
actual y	no	yes	Row Total

no	785	15	800
	0.981	0.019	0.884
	0.891	0.625	
	0.867	0.017	

yes	96	9	105
	0.914	0.086	0.116
	0.109	0.375	
	0.106	0.010	

Column Total	881	24	905
	0.973	0.027	

Cell Contents

N
N / Row Total
N / Col Total
N / Table Total

Total Observations in Table: 905

actual y	predicted y		Row Total
	no	yes	
no	787	13	800
	0.984	0.016	0.884
	0.896	0.481	
	0.870	0.014	
yes	91	14	105
	0.867	0.133	0.116
	0.104	0.519	
	0.101	0.015	
Column Total	878	27	905
	0.970	0.030	

The kappa for trials 10, 20, 30: 0.869, 0.864, 0.873

Error rate followed by increasing order of trials:

```
[1] 0.1182320 0.1226519 0.1149171
```

Trials 30 has a better result. But after running 40, 50, 60 trials (too slow the process to put it here in the report), trials 30 is still the best approach, while it remains similar results with decision tree model.

****I started to think if the results would be better with evenly distributed split partition. So I created another basic decision tree model using createDataPartition(). Check the proportion of train and test data:**

```
no    yes
0.8847111 0.1152889
```

```
no    yes
0.8849558 0.1150442
```

Great. Then use C5.0 to generate a tree:

Call:

```
C5.0.default(x = f_train_i[-16], y = f_train_i$y)
```

Classification Tree

Number of samples: 3617

Number of predictors: 15

Tree size: 4

Non-standard options: attempt to group attributes

```
no yes  
868 36
```

After prediction, run a CrossTable() to see that the kappa is 0.892.

```
Cell Contents  
-----|  
          N |  
Chi-square contribution |  
N / Row Total |  
N / Col Total |  
N / Table Total |  
-----|
```

Total Observations in Table: 904

		predicted y		
actual y		no	yes	Row Total
no	-----		-----	-----
	no	789	11	800
		0.566	13.656	
		0.986	0.014	0.885
		0.909	0.306	
yes	-----		-----	-----
	yes	79	25	104
		4.357	105.050	
		0.760	0.240	0.115
		0.091	0.694	
-----		-----	-----	-----
Column Total		868	36	904
		0.960	0.040	
-----		-----	-----	-----

Error rate:

```
[1] 0.09955752
```

Thus, this method has a better prediction result. I tried run boosting on that, failing to get a better result. I feel it too redundant to write it in the report, as different trials actually give the same results under such circumstances.

****Use ipred bagging to improve performance with nbagg = 25:**

```
f_pred_bag no yes
no 4000 11
yes 0 510
```

Error:

```
[1] 0.00243309
```

Cell Contents	
	N
Chi-square contribution	
N / Row Total	
N / Col Total	
N / Table Total	

Total Observations in Table: 4521

		f\$y		
f_pred_bag		no	yes	Row Total
no		4000	11	4011
		57.374	440.489	
		0.997	0.003	0.887
		1.000	0.021	
		0.885	0.002	
yes		0	510	510
		451.228	3464.319	
		0.000	1.000	0.113
		0.000	0.979	
		0.000	0.113	
Column Total		4000	521	4521
		0.885	0.115	

Kappa of CrossTable: 0.99

Both low error rate and high Kappa show that the Bagging method improved the prediction a lot. It is now the best approach.

****Use tuning in decision tree. Below is the summary of the tuning process.**

Bagged CART

```
4521 samples
15 predictor
2 classes: 'no', 'yes'
```

```
No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 4069, 4069, 4069, 4068, 4069, 4069, ...
Resampling results:
```

Accuracy Kappa
0.8845396 0.2238534

Call:
(function (x, y, trials = 1, rules = FALSE, weights = NULL, control = C5.0Control(), costs = NULL, ...) {
 1210L), .Names = c("subset", "bands", "winnow", "noGlobalPruning", "CF", "minCases",
 "fuzzyThreshold",
 "sample", "earlyStopping", "label", "seed")))

C5.0 [Release 2.07 GPL Edition] Sat Nov 4 23:37:01 2017

Class specified by attribute `outcome`

Read 4521 cases (42 attributes) from undefined.data

3 attributes winnowed

Estimated importance of remaining attributes:

- 44% contactunknown
- 7% poutcomesuccess
- 2% balance
- 1% jobtechnician
- 1% day
- 1% jobmanagement
- 1% maritalsingle
- 1% monthsep
- <1% maritalmarried
- <1% loanyes
- <1% monthaug
- <1% previous
- <1% age
- <1% jobblue-collar
- <1% jobentrepreneur
- <1% jobhousemaid
- <1% jobretired
- <1% jobself-employed
- <1% jobservices
- <1% jobstudent
- <1% jobunemployed
- <1% educationsecondary
- <1% educationtertiary
- <1% educationunknown
- <1% defaultyes
- <1% housingyes
- <1% contacttelephone
- <1% monthfeb
- <1% monthjan
- <1% monthjul
- <1% monthjun
- <1% monthmay
- <1% monthnov
- <1% monthoct
- <1% campaign

<1% pdays
<1% poutcomeother
<1% poutcomeunknown

----- Trial 0: -----

Rules:

Rule 0/1: (4392/438, lift 1.0)
poutcomesuccess <= 0
-> class no [0.900]

Rule 0/2: (129/46, lift 5.6)
poutcomesuccess > 0
-> class yes [0.641]

Default class: no

----- Trial 1: -----

Rules:

Rule 1/1: (4178.3/1082.7, lift 1.0)
age <= 62
monthoct <= 0
pdays <= 374
-> class no [0.741]

Rule 1/2: (49.3/13, lift 2.5)
monthoct <= 0
pdays > 374
-> class yes [0.728]

Rule 1/3: (132.3/41.8, lift 2.4)
monthoct > 0
-> class yes [0.682]

Rule 1/4: (190.8/75.2, lift 2.1)
age > 62
pdays <= 374
-> class yes [0.605]

Default class: no

----- Trial 2: -----

Rules:

Rule 2/1: (1068.7/228.3, lift 1.3)
contactunknown > 0
-> class no [0.786]

Rule 2/2: (3908.3/1426, lift 1.0)
jobstudent <= 0
monthjun <= 0
-> class no [0.635]

Rule 2/3: (147.3/34, lift 2.1)

contactunknown <= 0
monthjun > 0
-> class yes [0.765]

Rule 2/4: (95.1/33.8, lift 1.7)
jobstudent > 0
contactunknown <= 0
monthjun <= 0
-> class yes [0.641]

Default class: no

----- Trial 3: -----

Rules:

Rule 3/1: (4241.4/1775.5, lift 1.0)
poutcomesuccess <= 0
-> class no [0.581]

Rule 3/2: (279.6/110.2, lift 1.4)
poutcomesuccess > 0
-> class yes [0.605]

Default class: no

----- Trial 4: -----

Rules:

Rule 4/1: (979.8/289, lift 1.3)
contactunknown > 0
-> class no [0.705]

Rule 4/2: (542.5/185, lift 1.2)
loanyes > 0
-> class no [0.658]

Rule 4/3: (618.3/229.5, lift 1.2)
balance <= 1493
loanyes <= 0
-> class no [0.628]

Rule 4/4: (2675.8/1196.9, lift 1.2)
balance > 1493
loanyes <= 0
contactunknown <= 0
-> class yes [0.553]

Default class: no

----- Trial 5: -----

Rules:

Rule 5/1: (124/27.2, lift 1.4)
balance > 9122
monthjun <= 0

monthsep <= 0
-> class no [0.776]

Rule 5/2: (930.4/295.9, lift 1.2)
contactunknown > 0
-> class no [0.682]

Rule 5/3: (2977/1232.7, lift 1.0)
balance <= 3334
monthjun <= 0
monthsep <= 0
pdays <= 404
-> class no [0.586]

Rule 5/4: (46.8/10.5, lift 1.7)
contactunknown <= 0
monthjun <= 0
monthsep <= 0
pdays > 404
-> class yes [0.764]

Rule 5/5: (164.4/60.9, lift 1.4)
contactunknown <= 0
monthjun > 0
-> class yes [0.628]

Rule 5/6: (81.6/32.4, lift 1.4)
contactunknown <= 0
monthsep > 0
-> class yes [0.600]

Rule 5/7: (761.3/315.2, lift 1.3)
balance > 3334
balance <= 9122
contactunknown <= 0
monthsep <= 0
-> class yes [0.586]

Default class: no

----- Trial 6: -----

Rules:

Rule 6/1: (104.9/23.2, lift 1.4)
educationunknown <= 0
campaign > 9
-> class no [0.774]

Rule 6/2: (155.3/43.6, lift 1.3)
educationunknown > 0
monthoct <= 0
-> class no [0.716]

Rule 6/3: (2072.1/748, lift 1.2)
maritalmarried > 0
monthoct <= 0
previous <= 1

-> class no [0.639]

Rule 6/4: (177.8/71, lift 1.3)
monthoct > 0
-> class yes [0.599]

Rule 6/5: (675.7/285.5, lift 1.3)
educationunknown <= 0
monthoct <= 0
campaign <= 9
previous > 1
-> class yes [0.577]

Rule 6/6: (1821.2/851.8, lift 1.2)
maritalmarried <= 0
educationunknown <= 0
campaign <= 9
-> class yes [0.532]

Default class: no

----- Trial 7: -----

Rules:

Default class: no

*** boosting reduced to 7 trials since last classifier is very inaccurate

Evaluation on training data (4521 cases):

Trial	Rules
-----	-----
	No Errors
0	2 484(10.7%)
1	4 555(12.3%)
2	4 565(12.5%)
3	2 484(10.7%)
4	4 1965(43.5%)
5	7 850(18.8%)
6	6 1790(39.6%)
boost	492(10.9%) <<

(a) (b) <-classified as

-----	-----	
3927	73	(a): class no
419	102	(b): class yes

Attribute usage:

100.00%	poutcomesuccess
99.65%	pdays
99.62%	monthoct
97.97%	age

90.02%	monthjun
89.38%	maritalmarried
88.52%	balance
88.08%	jobstudent
85.80%	monthsep
80.42%	contactunknown
79.81%	loanyes
64.83%	previous
49.10%	educationunknown
45.08%	campaign

Time: 0.2 secs

```
f_pred_tu no yes
no 3927 419
yes 73 102
```

Cell Contents

N
Chi-square contribution
N / Row Total
N / Col Total
N / Table Total

Total Observations in Table: 4521

f_pred_tu	no	yes	Row Total
no	3927	419	4346
	1.742	13.371	
	0.904	0.096	0.961
	0.982	0.804	
	0.869	0.093	
yes	73	102	175
	43.251	332.059	
	0.417	0.583	0.039
	0.018	0.196	
	0.016	0.023	
Column Total	4000	521	4521
	0.885	0.115	

CrossTable show that the $[Pr(a)-Pr(e)]/[1-Pr(3)] = 0.89$, pretty good, but not as good as ipred bagging.

**Use customized bagging, and below is the result:

```
.model .trials .winnnow
1 tree 1 FALSE
2 tree 5 FALSE
3 tree 10 FALSE
4 tree 15 FALSE
5 tree 20 FALSE
6 tree 25 FALSE
7 tree 30 FALSE
8 tree 35 FALSE
C5.0
```

```
4521 samples
15 predictor
2 classes: 'no', 'yes'
```

No pre-processing
 Resampling: Cross-Validated (10 fold)
 Summary of sample sizes: 4069, 4069, 4069, 4068, 4069, 4069, ...
 Resampling results across tuning parameters:

trials	Accuracy	Kappa
1	0.8922801	0.2325560
5	0.8854236	0.1404052
10	0.8852004	0.1625564
15	0.8829885	0.1749768
20	0.8854222	0.1753795
25	0.8858646	0.1874779
30	0.8856434	0.1868124
35	0.8856434	0.1868124

Tuning parameter 'model' was held constant at a value of tree
 Tuning parameter 'winnnow' was held constant at a value of FALSE
 Kappa was used to select the optimal model using the one SE rule.
 The final values used for the model were trials = 1, model = tree and winnow = FALSE.

It shows that when trial = 1, the model has the best prediction, but not as good as ipred bagging.

Step 4.5 Another Model

Use Random Forest

```
Call:
randomForest(formula = y ~ ., data = f)
Type of random forest: classification
Number of trees: 500
No. of variables tried at each split: 3
```

```
OOB estimate of error rate: 11.1%
Confusion matrix:
no yes class.error
no 3934 66 0.0165000
yes 436 85 0.8368522
```

I built a random forest on the data, but the result doesn't look good. It appeared to be much worse than most of the models above. Thus, I started resampling using both random forest and C50. The results are as follows, in the order of Random Forest resampling and C50 resampling:

Random Forest

4521 samples
15 predictor
2 classes: 'no', 'yes'

No pre-processing

Resampling: Cross-Validated (10 fold, repeated 10 times)

Summary of sample sizes: 4069, 4069, 4069, 4069, 4068, 4069, ...

Resampling results across tuning parameters:

mtry	Accuracy	Kappa
2	0.8854902	0.02817444
4	0.8900023	0.15028977
8	0.8887418	0.18531178
16	0.8885428	0.21570099

Kappa was used to select the optimal model using the largest value.

The final value used for the model was mtry = 16.

C5.0

4521 samples
15 predictor
2 classes: 'no', 'yes'

No pre-processing

Resampling: Cross-Validated (10 fold, repeated 10 times)

Summary of sample sizes: 4069, 4069, 4069, 4068, 4069, 4069, ...

Resampling results across tuning parameters:

trials	Accuracy	Kappa
10	0.8879229	0.1729337
20	0.8880113	0.2046911
30	0.8877680	0.2035707

Tuning parameter 'model' was held constant at a value of tree

Tuning parameter 'winnow' was held constant at a

value of FALSE

Kappa was used to select the optimal model using the largest value.

The final values used for the model were trials = 20, model = tree and winnow = FALSE.

Step 5 - Select the final model

During my work, I found out that most models have accuracy in 85% to 90%, which are all pretty decent results. In pragmatic world, these figures could give a good predictions on future trends. However, in decision tree model with ipred bagging, the accuracy is over 99%. This is quite impressive that I tried to run it more times to make sure no bug exists. In the end, I would pick this method as the best approach.

In addition to conclusion, I had some interesting discoveries. In common sense, more trials in machine learning might lead to more accuracy, while it had been proved wrong in my work. Also, when boosting fail to improve the accuracy in one scenario, I should still try it in other models, because it might work.

Overall I am glad to have this final exam as it taught me a lot of insights during practice.